

# Building a Bigram Language Model from Scratch



---

This is an extended version of Andrej Karpathy's notebook in addition to his [Zero To Hero](#) video on bigram language models.

Adapted by:

Prof. Dr.-Ing. Antje Muntzinger, University of Applied Sciences Stuttgart

[antje.muntzinger@hft-stuttgart.de](mailto:antje.muntzinger@hft-stuttgart.de)

---

We'll construct a bigram language model from scratch. The model will be trained on a text file containing names and will be able to generate new names based on what it has learned. So if you are expecting a baby and looking for some extraordinary new name, you might get some inspiration here :)

## Table of Contents

1. Loading the Data
2. Preprocessing the Data
3. Bigram Counts as PyTorch Tensor
4. Sampling New Characters
5. Broadcasting
6. Generating New Names
7. Evaluating the Model
8. The Neural Network Approach
9. The First Neuron
10. Creating 27 Neurons
11. Summary (so far)
12. Optimization
13. Optimizing over the Whole Dataset

## 14. Sampling from the Neural Net

## 15. Conclusion

# 1. Loading and Preprocessing the Data

First, we read the names from the text file and save them in a list.

```
In [1]: words = open('names.txt', 'r').read().splitlines()
```

**TODD:** 1) Print the first 10 names. Find out the number of names contained in the dataset as well as the shortest and longest name. **(2 points)**

```
In [2]: # YOUR CODE GOES HERE
print("First 10 names in the dataset: ", words[:10])
print("Number of names contained in the dataset: ", len(words))
print("Longest name in the dataset: ", max(words, key=len))
print("Shortest name in the dataset: ", min(words, key=len))
```

First 10 names in the dataset: ['emma', 'olivia', 'ava', 'isabella', 'sophia', 'charlotte', 'mia', 'amelia', 'harper', 'evelyn']  
Number of names contained in the dataset: 32033  
Longest name in the dataset: muhammadibrahim  
Shortest name in the dataset: an



Note that in one example name like "isabella", a lot of information is contained. For example:

- "i" is likely to be the first character of a name
- "s" is likely to follow after an "i"
- "a" is likely to follow after "is"
- ...
- after "isabella", the name is likely to end.

Here we are going to create a **bigram** language model, which means we only use the single previous character to predict the next. For example, we use "s" (not "is") to predict "a" in "isabella" and forget that we have a lot more information.

# 2. Preprocessing the Data

First, we introduce two special characters, `<S>` for "start" and `<E>` for "end". Then we zip two consecutive characters and print them:

```
In [3]: b = {} # dictionary to store bigram counts
for w in words: # print bigrams
    chs = ['<S>'] + list(w) + ['<E>']
    for ch1, ch2 in zip(chs, chs[1:]): # zip(['a', 'b', 'c'], ['b', 'c', 'd']) -> [
        bigram = (ch1, ch2)
        # count how many times bigram appears
```

```
b[bigram] = b.get(bigram, 0) + 1 # b.get(bigram, 0) returns b[bigram] if bi
print(ch1, ch2, b[bigram])
```

```
<S> e 1
e m 1
m m 1
m a 1
a <E> 1
<S> o 1
o l 1
l i 1
i v 1
v i 1
i a 1
a <E> 2
<S> a 1
a v 1
v a 1
a <E> 3
<S> i 1
i s 1
s a 1
a b 1
...truncated...
```

We see that already in the first three names, the bigram `a <E>` occurs three times - which makes sense because many names end with an "a".

```
In [4]: # print all bigrams
b
```

```
In [5]: # sort by the count (by default it sorts by the first element of the tuple, but we
sorted(b.items(), key = lambda kv: kv[1], reverse=True) # kv is a tuple, kv[0] is t
```

**TODO:** 2) What is the most common bigram in the dataset? What is the least common bigram? **(2 points)**

```
In [6]: # YOUR CODE GOES HERE
print("most common bigram:", max(b, key=b.get))
print("least common bigram:", min(b, key=b.get))
```

```
most common bigram: ('n', '<E>')
least common bigram: ('q', 'r')
```



### 3. Bigram Counts as PyTorch Tensor

Next, we store the bigrams in a matrix instead of a python dictionary: The rows are going to be the first character of the bigram and the columns are going to be the second character. Each entry in this matrix will tell us how often that first character is followed by the second character in the dataset.

To create this matrix (also called **2D array** or **2D tensor**), we use PyTorch, which allows us to create multi-dimensional arrays and manipulate them very efficiently. We will also use a special character `.` for word boundaries instead of `<A>` and `<E>` because we don't need to distinguish between these two.

```
In [7]: import torch
```

We get all characters and map them to integers in a mapping called "s to i":

```
In [8]: chars = sorted(list(set(''.join(words)))) # get all unique characters in words
stoi = {s:i+1 for i,s in enumerate(chars)} # create a dictionary mapping characters
stoi['.'] = 0 # add a special character for word boundaries instead of <A> and <E>
```


We also create the inverse mapping "i to s" that assigns a character to each integer. We can use `items()` to get a list of a dictionary's tuples:

```
In [9]: print(stoi)
print(stoi.items())
```

```
{ 'a': 1, 'b': 2, 'c': 3, 'd': 4, 'e': 5, 'f': 6, 'g': 7, 'h': 8, 'i': 9, 'j': 10,
'k': 11, 'l': 12, 'm': 13, 'n': 14, 'o': 15, 'p': 16, 'q': 17, 'r': 18, 's': 19,
't': 20, 'u': 21, 'v': 22, 'w': 23, 'x': 24, 'y': 25, 'z': 26, '.': 0}
dict_items([('a', 1), ('b', 2), ('c', 3), ('d', 4), ('e', 5), ('f', 6), ('g', 7),
('h', 8), ('i', 9), ('j', 10), ('k', 11), ('l', 12), ('m', 13), ('n', 14), ('o', 15),
('p', 16), ('q', 17), ('r', 18), ('s', 19), ('t', 20), ('u', 21), ('v', 22),
('w', 23), ('x', 24), ('y', 25), ('z', 26), ('.', 0)])
```

```
In [10]: itos = {i:s for s,i in stoi.items()} # reverse mapping from integers to characters
```

**TODO:** 3a) We now have a vocabulary of 27 characters. Instantiate a torch tensor called `N` of size 27x27 with zero values of dtype int32. Check the PyTorch documentation to see how to create a tensor if needed. **(1 point)**

```
In [11]: # YOUR CODE GOES HERE
N = torch.zeros((len(stoi), len(stoi)), dtype=torch.int32) 
```

Let's fill the matrix `N` with the counts of the bigrams:

```
In [12]: # store the counts of bigrams in the matrix
for w in words:
    chs = ['.'] + list(w) + ['.']
    for ch1, ch2 in zip(chs, chs[1:]):
        ix1 = stoi[ch1]
        ix2 = stoi[ch2]
        N[ix1, ix2] += 1
```

```
In [13]: # print the matrix
import matplotlib.pyplot as plt
%matplotlib inline

plt.figure(figsize=(16,16))
```

```
plt.imshow(N, cmap='Blues')
for i in range(27):
    for j in range(27):
        chstr = itos[i] + itos[j] # character string of bigram, e.g. 'ab'
        plt.text(j, i, chstr, ha="center", va="bottom", color='gray') # upper text:
        plt.text(j, i, N[i, j].item(), ha="center", va="top", color='gray') # Lower
plt.axis('off')
```

Out[13]: (-0.5, 26.5, 26.5, -0.5)

.o	.a	.b	.c	.d	.e	.f	.g	.h	.i	.j	.k	.l	.m	.n	.o	.p	.q	.r	.s	.t	.u	.v	.w	.x	.y	.z
0	4036	1306	1542	1690	1531	417	669	874	591	2422	2963	1572	2538	1146	394	515	92	1639	2055	1308	78	376	307	134	535	929
a.6640	aa556	ab541	ac470	ad1042	ae692	af134	ag168	ah2332	ai1650	aj175	ak568	al2528	am1634	an5438	ao63	ap82	aq60	ar3264	as1118	at687	au381	av834	aw161	ax182	ay2050	az435
b.114	ba321	bb38	bc1	bd65	be655	bf0	bg0	bh41	bi217	bj1	bk0	bl103	bm0	bn4	bo105	bp0	bq0	br842	bs8	bt2	bu45	bv0	bw0	bx0	by83	bz0
c.97	ca815	cb0	cc42	cd1	ce551	cf0	cg2	ch664	ci271	cj3	ck316	cl116	cm0	cn0	co380	cp1	cq11	cr76	cs5	ct35	cu35	cv0	cw0	cx3	cy104	cz4
d.516	da1303	db1	dc3	dd149	de1283	df5	dg25	dh118	di674	dj9	dk3	dl60	dm30	dn31	do378	dp0	dq1	dr424	ds29	dt4	du92	dv17	dw23	dx0	dy317	dz1
e.5952	ea679	eb121	ec153	ed384	ee1271	ef82	eg125	eh152	ei818	ej55	ek178	el3248	em769	en2675	eo269	ep83	eq14	er1958	es861	et580	eu69	ev463	ew50	ex132	ey1070	ez181
f.80	fa242	fb0	fc0	fd0	fe123	ff44	fg1	fh1	fi160	fj0	fk2	fl20	fm0	fn4	fo60	fp0	fq0	fr114	fs6	ft18	fu10	fv0	fw4	fx0	fy14	fz2
g.108	ga330	gb3	gc0	gd19	ge334	gf1	gg25	gh360	gi190	gj3	gk0	gl32	gm6	gn27	go83	gp0	gq0	gr201	gs30	gt31	gu85	gv1	gw26	gx0	gy31	gz1
h.2409	ha2244	hb8	hc2	hd24	he674	hf2	hg2	hh1	hi729	hj9	hk29	hl185	hm117	hn138	ho287	hp1	hq1	hr204	hs31	ht71	hu166	hv39	hw10	hx0	hy213	hz20
i.2489	ia2445	ib110	ic509	id440	ie1653	if101	ig148	ih95	ii82	ij76	ik445	il1345	im427	in2126	io588	ip53	iq52	ir849	is1316	it541	iu109	iv269	iw8	ix89	iy779	iz277
j.71	ja1473	jb1	jc4	jd4	je440	jf0	fg0	jh45	ji119	jj2	jk2	jl9	jm5	jn2	jo479	jp1	jq0	jr11	js7	jt2	ju202	jv5	jw6	jx0	jy10	jz0
k.363	ka1731	kb2	kc2	kd2	ke895	kf1	kg0	kh307	ki509	kj2	kk20	kl139	km9	kn26	ko344	kp0	kq0	kr109	ks95	kt17	ku50	kv2	kw34	kx0	ky379	kz2
l.1314	la2623	lb52	lc25	ld138	le2921	lf22	lg6	lh19	li2480	lj6	lk24	ll1345	lm60	ln14	lo692	lp15	lq3	lr18	ls94	lt77	lu324	lv72	lw16	lx0	ly1588	lz10
m.516	ma2590	mb112	mc51	md24	me818	mf1	mg0	mh5	mi1256	mj7	mk1	ml5	mm168	mn20	mo452	mp38	mq0	mr97	ms35	mt4	mu139	mv3	mw2	mx0	my287	mz11
n.6763	na2977	nb8	nc213	nd704	ne1359	nf11	ng273	nh26	ni1725	nj44	nk58	nl195	nm19	nn1906	no496	np5	nq2	nr44	ns278	nt443	nu96	nv55	nw11	nx6	ny465	nz145
o.855	oa149	ob140	oc114	od190	oe132	of34	og44	oh171	oi69	oj16	ok68	ol619	om261	on2411	oo115	op95	oq3	or1059	os504	ot118	ou275	ov176	ow114	ox45	oy103	oz54
p.33	pa209	pb2	pc1	pd0	pe197	pf1	pg0	ph204	pi61	pj1	pk1	pl16	pm1	pn1	po59	pp39	pq0	pr151	ps16	pt17	pu4	pv0	pw0	px0	py12	pz0
q.28	qa13	qb0	qc0	qd0	qe1	qf0	qg0	qh0	qi13	qj0	qk0	ql1	qm2	qn0	qo2	qp0	qq0	qr1	qs2	qt0	qu206	qv0	qw3	qx0	qy0	qz0
r.1377	ra2356	rb41	rc99	rd187	re1697	rf9	rg76	rh121	ri3033	rj25	rk90	rl413	rm162	rn140	ro869	rp14	rq16	rr425	rs190	rt208	ru252	rv80	rw21	rx3	ry773	rz23
s.1169	sa1201	sb21	sc60	sd9	se884	sf2	sg2	sh1285	si684	sj2	sk82	sl279	sm90	sn24	so531	sp51	sq1	sr55	ss461	st765	su185	sv14	sw24	sx0	sy215	sz10
t.483	ta1027	tb1	tc17	td0	te716	tf2	tg2	th647	ti532	tj3	tk0	tl134	tm4	tn22	to667	tp0	tq0	tr352	ts35	tt374	tu78	tv15	tw11	tx2	ty341	tz105
u.155	ua163	ub103	uc103	ud136	ue169	uf19	ug47	uh58	ui121	uj14	uk93	ul301	um154	un275	uo10	up16	uq10	ur414	us474	ut82	uu3	uv37	uw86	ux34	uy13	uz45
v.88	va642	vb1	vc0	vd1	ve568	vf0	vg0	vh1	vi911	vj0	vk3	vl14	vm0	vn8	vo153	vp0	vq0	vr48	vs0	vt0	vu7	vv7	vw0	vx0	vy121	vz0
w.51	wa280	wb1	wc0	wd8	we149	wf2	wg1	wh23	wi148	wj0	wk6	wl13	wm2	wn58	wo36	wp0	wq0	wr22	ws20	wt8	wu25	wv0	ww2	wx0	wy73	wz1
x.164	xa103	xb1	xc4	xd5	xe36	xf3	xg0	xh1	xi102	xj0	xk0	xl39	xm1	xn1	xo41	xp0	xq0	xr0	xs31	xt70	xu5	xv0	xw3	xx38	xy30	xz19
y.2007	ya2143	yb27	yc115	yd272	ye301	yf12	yg30	yh22	yi192	yj23	yk86	yl1104	ym148	yn1826	yo271	yp15	yq6	yr291	ys401	yt104	yu141	yv106	yw4	yx28	yy23	yz78
z.160	za860	zb4	zc2	zd2	ze373	zf0	zg1	zh43	zi364	zj2	zk2	zl123	zm35	zn4	zo110	zp2	zq0	zr32	zs4	zt4	zu73	zv2	zw3	zx1	zy147	zz45

**TODO:** 3b) Why do we use `.item()` in the code cell above? Check out the PyTorch documentation if needed! (1 point)

**ANSWER:** YOUR ANSWER GOES HERE

The `.item()` method is used to extract the value of a single element from a PyTorch tensor and convert it from the tensor `N[i, j]` (which is a 0-dimensional tensor) to an integer before

being passed to `plt.text()`. This is necessary because `plt.text()` expects standard Python data types, not PyTorch tensor objects.

Without `.item()`, PyTorch would return a tensor object (like `tensor(5)`), which might cause issues or unexpected behavior when trying to display the value using matplotlib.

**TODO:** 3c) Assume our first character is an 'f'. Looking at the matrix plot above, which is the most likely next character? **(1 point)**

```
In [14]: # YOUR ANSWER GOES HERE
print(itos[torch.argmax(N[stoi['f']]).item()])
```

a

```
In [15]: print(N[1,2])
print(N[1,2].item())
```

```
tensor(541, dtype=torch.int32)
541
```

Here is a visual summary so far - this is how we can get the most likely next character for input 'e' by plucking out the 6th row:



No description has been provided for this image

## 4. Sampling New Characters

Now we want to sample new characters. We start with the first character `.`, and then sample the next character based on the count of bigrams. We have the raw counts stored in the matrix `N`, but we need to convert them to probabilities in order to sample from the distribution.

**TODO:** 4a) Pluck out the first row of the matrix, convert the entries to floats, and normalize them so that the row sums to 1. **(3 points)**

**HINT:** You can simply cast a PyTorch tensor to float using `.float()`, see for example here: <https://www.datascienceweekly.org/tutorials/pytorch-change-tensor-type-cast-a-pytorch-tensor-to-another-type>

```
In [16]: # YOUR CODE GOES HERE
p = N[0].float() # convert first row of N to float
p /= p.sum() # normalize values
print(p)
```

```
tensor([0.0000, 0.1377, 0.0408, 0.0481, 0.0528, 0.0478, 0.0130, 0.0209, 0.0273,
        0.0184, 0.0756, 0.0925, 0.0491, 0.0792, 0.0358, 0.0123, 0.0161, 0.0029,
        0.0512, 0.0642, 0.0408, 0.0024, 0.0117, 0.0096, 0.0042, 0.0167, 0.0290])
```

**TODO:** 4b) How can you interpret this row now? **(1 point)**

**ANSWER:** YOUR ANSWER GOES HERE

All values are between 0 and 1 and can be interpreted as a nominal probability.

The highest possibility is at 13.77% for .a whereas the lowest possibility lies at 0.24% for .u

To sample from this distribution, we use `torch.multinomial`, which samples from the multinomial probability distribution (in simple words: "you give me probabilities and I will give you integers sampled according to the probability distribution").

```
In [17]: g = torch.Generator().manual_seed(2147483647) # makes the random numbers reproducib
ix = torch.multinomial(p, num_samples=1, replacement=True, generator=g).item() # sa
# replacement=True means that we are sampling with replacement (i.e. we can sample
print("sampled index:", ix) # print the sampled index
itos[ix] # convert the sampled index to a character
```

sampled index: 10

Out[17]: 'j'

**TODO:** 4c) Sample 100 characters according to your probability distribution for word beginnings. Print the sampled indices and characters. Can you relate the output to your probability distribution? Did you expect this output, or can you see unexpected output characters? **(2 points)**

```
In [18]: # YOUR CODE GOES HERE
g = torch.Generator().manual_seed(2147483647)
ix_100 = torch.multinomial(first_row, num_samples=100, replacement=True, generator=
for i in ix_100:
    print(i.item(), ":", itos[i.item()])
```


```
-----
NameError                                Traceback (most recent call last)
Cell In[18], line 3
      1 # YOUR CODE GOES HERE
      2 g = torch.Generator().manual_seed(2147483647)
----> 3 ix_100 = torch.multinomial(first_row, num_samples=100, replacement=True, gen
erator=g) # use torch.multinomial with batch size of 100
      4 for i in ix_100:
      5     print(i.item(), ":", itos[i.item()])

NameError: name 'first_row' is not defined
```

**ANSWER:** YOUR ANSWER GOES HERE

Yes, i can relate the output to probability distribution. I expect this output because the letters with highest probabilities are : a,b,c,d,e,j,k,l,m,n and they are found in output more than others. Also the probability of '.' was 0 and i have never seen it on the output.

This is the updated visualization including normalization for probabilities as outputs:

No description has been provided for this image

## 5. Broadcasting

We store all probabilities in a matrix `P`, so that each row is normalized to 1 and contains the probabilities for the next character:


```
In [19]: P = (N+1).float()
P /= P.sum(1, keepdims=True) # normalize the rows of P (we use keepdims=True to keep
print(P)
```

```
tensor([[3.1192e-05, 1.3759e-01, 4.0767e-02, 4.8129e-02, 5.2745e-02, 4.7785e-02,
        1.3038e-02, 2.0898e-02, 2.7293e-02, 1.8465e-02, 7.5577e-02, 9.2452e-02,
        4.9064e-02, 7.9195e-02, 3.5777e-02, 1.2321e-02, 1.6095e-02, 2.9008e-03,
        5.1154e-02, 6.4130e-02, 4.0830e-02, 2.4641e-03, 1.1759e-02, 9.6070e-03,
        4.2109e-03, 1.6719e-02, 2.9008e-02],
        [1.9583e-01, 1.6425e-02, 1.5983e-02, 1.3889e-02, 3.0756e-02, 2.0435e-02,
        3.9809e-03, 4.9835e-03, 6.8796e-02, 4.8685e-02, 5.1899e-03, 1.6779e-02,
        7.4575e-02, 4.8213e-02, 1.6039e-01, 1.8872e-03, 2.4475e-03, 1.7988e-03,
        9.6279e-02, 3.2997e-02, 2.0288e-02, 1.1264e-02, 2.4623e-02, 4.7771e-03,
        5.3963e-03, 6.0480e-02, 1.2857e-02],
        [4.3039e-02, 1.2051e-01, 1.4596e-02, 7.4850e-04, 2.4701e-02, 2.4551e-01,
        3.7425e-04, 3.7425e-04, 1.5719e-02, 8.1587e-02, 7.4850e-04, 3.7425e-04,
        3.8922e-02, 3.7425e-04, 1.8713e-03, 3.9671e-02, 3.7425e-04, 3.7425e-04,
        3.1549e-01, 3.3683e-03, 1.1228e-03, 1.7216e-02, 3.7425e-04, 3.7425e-04,
        3.7425e-04, 3.1437e-02, 3.7425e-04],
        [2.7536e-02, 2.2928e-01, 2.8098e-04, 1.2082e-02, 5.6196e-04, 1.5510e-01,
        2.8098e-04, 8.4293e-04, 1.8685e-01, 7.6426e-02, 1.1239e-03, 8.9070e-02,
        3.2874e-02, 2.8098e-04, 2.8098e-04, 1.0705e-01, 5.6196e-04, 3.3717e-03,
        2.1635e-02, 1.6859e-03, 1.0115e-02, 1.0115e-02, 2.8098e-04, 2.8098e-04,
        1.1239e-03, 2.9503e-02, 1.4049e-03],
        ...truncated...]
```

**TODO:** 5) Can you think of a reason why we calculate the probabilities based on `N+1` instead of `N`? (1 point)

**HINT:** What happens for `N=0`?

**ANSWER:** YOUR ANSWER GOES HERE

It is a so called Laplace-smoothing which makes sure that no entry is 0. This is needed not have any 0% probability for the entries without occurrence in the training data. 

We need `keepdims=True` in order to sum each line, the result should be a 27x1 column vector containing the row sums - otherwise all entries get collapsed to a 1D instead of 2D array:



```
In [20]: print(P.sum(1, keepdims=True).shape) # shape collapses to 27x1 due to summation along dim 1
print(P.sum(1, keepdims=False).shape) # 1D tensor of size 27

torch.Size([27, 1])
torch.Size([27])
```

But why does the division above `P /= P.sum(1, keepdims=True)` even work? We divide the 27x27 matrix `P` by a 27x1 vector. This only works because PyTorch automatically applies **broadcasting**, a type of tensor manipulation: From <https://pytorch.org/docs/stable/notes/broadcasting.html> we get the information that two tensors are **broadcastable** if the following rules hold:

Each tensor has at least one dimension.

When iterating over the dimension sizes, starting at the trailing dimension, the dimension sizes must either be equal, one of them is 1, or one of them does not exist.

If two tensors `x, y` are "broadcastable", the resulting tensor size is calculated as follows:

If the number of dimensions of `x` and `y` are not equal, prepend 1 to the dimensions of the tensor with fewer dimensions to make them equal length.

Then, for each dimension size, the resulting dimension size is the max of the sizes of `x` and `y` along that dimension.

In our case, broadcasting takes the 27x1 column vector and copies it 27 times to get a 27x27 matrix, then it makes an element-wise division. We can check the result:

```
In [21]: # check that the first row sums to 1
print(P[0].sum())
```

```
tensor(1.)
```

**TODO (optional):** What happens if you remove `keepdims=True` ? Will the division still result in the same matrix - why or why not?

**HINT:** Don't forget to revert your changes for later cells to work right - or copy the code cell and use different variable names.

**ANSWER:** YOUR ANSWER GOES HERE

If we remove it, it will return 1D array(each element is sum of one row),then when we divide the tensor(27,27) with sum(27). As it didnt protect the dimension features, the division would be wrong(maybe column wise). In the code below, as there is a mistake in calculation, sum of first row is not 1 :



In [22]: *# YOUR CODE GOES HERE*

```
P_v2 = (N+1).float()
P_v2 /= P_v2.sum(1)
print("If it would normalized correctly sum supposed to be 1 but : ",P_v2[0].sum())

# Let us see in a small example
temp=torch.tensor([[1,2,3],[4,5,6],[7,8,9]]).float()
print("\nwithout keepdim=False :\n ",temp.sum(1, keepdim=False) , "\n normalized:\n"

print("\nwith keepdim=True :\n",temp.sum(1,keepdim=True) , "\n normalized:\n" , tem
```

If it would normalized correctly sum supposed to be 1 but : tensor(6.9246)

```
without keepdim=False :
    tensor([ 6., 15., 24.])
normalized:
tensor([[0.1667, 0.1333, 0.1250],
        [0.6667, 0.3333, 0.2500],
        [1.1667, 0.5333, 0.3750]])
```

```
with keepdim=True :
    tensor([[ 6.],
            [15.],
            [24.]])
normalized:
tensor([[0.1667, 0.3333, 0.5000],
        [0.2667, 0.3333, 0.4000],
        [0.2917, 0.3333, 0.3750]])
```

## 6. Generating New Names

With this, we can sample new names by starting with the start character `.` and sampling the next character, then the next next character and so on, until we sample the end character `.`:

```
In [23]: g = torch.Generator().manual_seed(2147483647)

for i in range(20): # generate 20 names

    out = []
    ix = 0 # start with the start-of-word character
    while True:
        p = P[ix] # get the probabilities of the next character by plucking the row
        ix = torch.multinomial(p, num_samples=1, replacement=True, generator=g).ite
        out.append(itos[ix]) # convert the sampled index to a character and append
        if ix == 0: # if we sampled the end-of-word character: break
            break
    print(''.join(out)) # print the generated name
```

```
junide.  
janasah.  
p.  
cony.  
a.  
nn.  
kohin.  
tolian.  
juee.  
ksahnaauranilevias.  
dedainrwieta.  
ssonielylarte.  
faveumerifontume.  
phynslenaruani.  
core.  
yaenon.  
ka.  
jabdinerimikimaynin.  
anaasn.  
ssorionsush.  
...truncated...
```

This doesn't look right... but it is! The bigram language model is simply not powerful enough to create more reasonable names. For example, it generated a name "a" twice, but it doesn't know that "a" is the first character here. All it knows is that "a" is a character in the name and how likely it is that the name ends here, without looking at previous characters.

**TODO:** 6) In the code cell above, experiment with replacing `p` with the uniform distribution, making everything equally likely. Can you see that the bigram model works better than pure chance? **(2 points)**

**HINT:** Don't forget to revert your changes for later cells to work right - or copy the code cell and use different variable names.

```
In [24]: g = torch.Generator().manual_seed(2147483647)

for i in range(20): # generate 20 names
    out = []
    ix = 0 # start with the start-of-word character

    while True:
        # Replace `p` with a uniform distribution (all probabilities equal)
        p = torch.ones_like(P[ix]) # Create a vector with equal probabilities for
        p /= p.sum() # Normalize so the sum of probabilities is 1 (uniform distrib

        # Sample the next character
        ix = torch.multinomial(p, num_samples=1, replacement=True, generator=g).ite
        out.append(itos[ix]) # Add the character to the output list

        # Break the loop if the end-of-word character `.` is sampled
        if ix == 0:
            break
```



```
print(''.join(out))
```

```
juwjdvdipkcqaz.  
p.  
cfqywocnzqfjiirltozcogsjgwzvudlhnpauyjbilevhajkdbduinrwibtlzsnjyievyvaftbzffvmumthyf  
odtumjrpftytszwjhrjagq.  
coreaysezocfkyjjabdywejfmofmwyfinwagaasnhsvfihofszxhddgosfmptpagicz.  
rjpiufmthdt.  
rkrrsru.  
iyumuyfy.  
mjekujcbkhvupwyhvpvhvccragr.  
wdkhwfdztta.  
mplyisbxlyhuuiqzavmpocbztqmimvyqwat.  
f.  
.  
ndxjxfpvsqrtikyzzsaloevgvvnundewkfmbjzqegruxiteaxchwtmurzsodridcdznojvaliypvrghvxtezt  
wguciqqvywhqelv.  
viosvhibdhnceukgmtmboscnbzoiwupnwnpipixtewbgsgyewfdacbfcxrvjypkmsbranmjrdsydotafvkd  
kbdepihzpwsqdab.  
vfuo lwbasrtugtbbiqbujfdtskceqjtcldnfujqlspgkltalmlokdmsl.  
fddmxjv.  
mfsgxmw.  
vdihkvngtojvrdsyqivcob.  
uziengogtjvnvqgfjtkqufrxfjlglykiiluohgnoiuwzylq.  
fsgircvmhtipagkxwvjypnsriadmufjnlkcicvatjvryzeljxkbrlrjsp.  
...truncated...
```

## 7. Evaluating the Model

We want to evaluate the quality of the model using a single number, the **training loss**. But which single number to use? Let's print the probabilities of the first bigrams to start with:

```
In [25]: for w in words[:3]: # print bigram probabilities for the first 3 words  
        chs = ['.'] + list(w) + ['.'] # add start and end of word characters  
        for ch1, ch2 in zip(chs, chs[1:]): # iterate over bigrams  
            ix1 = stoi[ch1] # get the integer representation of the characters  
            ix2 = stoi[ch2] # get the integer representation of the characters  
            prob = P[ix1, ix2] # get the probability of the bigram  
            print(f'{ch1}{ch2}: {prob:.4f}') # print the bigram and its probability
```

.e: 0.0478  
em: 0.0377  
mm: 0.0253  
ma: 0.3885  
a.: 0.1958  
.o: 0.0123  
ol: 0.0779  
li: 0.1774  
iv: 0.0152  
vi: 0.3508  
ia: 0.1380  
a.: 0.1958  
.a: 0.1376  
av: 0.0246  
va: 0.2473  
a.: 0.1958

**TODO:** 7a) What probabilities would we expect for an untrained model? How do you interpret these outputs compared to the untrained model? **(2 points)**

**ANSWER:** YOUR ANSWER GOES HERE

For an untrained model, we would expect all bigram probabilities to be equal. This is because, in an untrained model, we have no prior information about the likelihood of different character transitions.

If there are 27 possible characters (including the start and end markers), each bigram probability in an untrained model would be close to  $1/27 \approx 0.037$



So how can we summarize these probabilities in a single number? First of all, we can use the **likelihood**, which is the product of all these probabilities. The likelihood tells us which probability our model assigns to the whole dataset. So we want the likelihood to be as big as possible for a good model (**maximum likelihood estimation**).

**TODO:** 7b) Can you think of a problem when calculating the likelihood as defined above? **(1 point)**

**HINT:** What happens numerically if we multiply many tiny numbers?

**ANSWER:** YOUR ANSWER GOES HERE

Numerical underflow becomes a problem. Since probabilities are often small values (less than 1), multiplying a large number of them can result in an extremely small product, potentially too small for the computer to represent accurately. This can lead to the calculated likelihood being rounded to zero, even when it's non-zero, causing a loss of precision in computations.

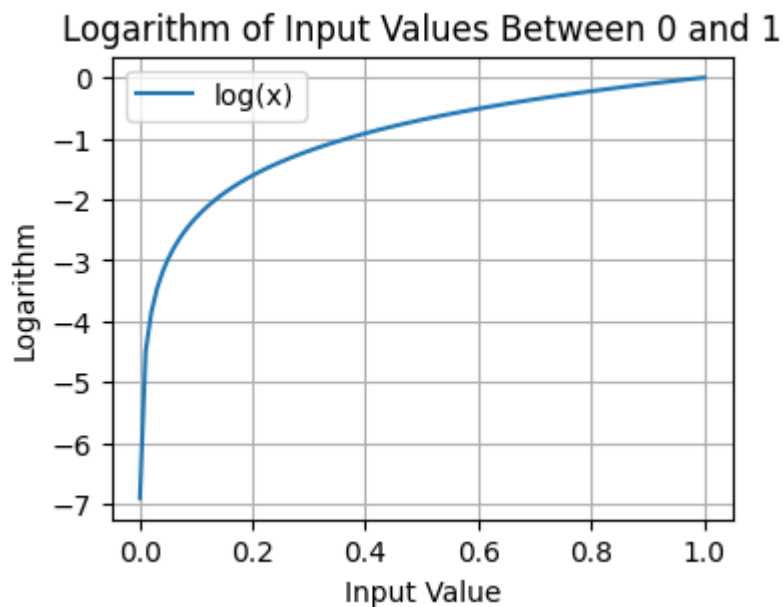


Instead of the likelihood, we use the **log likelihood**: Applying the logarithm transforms the likelihood to a value in  $[-\infty, 0]$ .

```
In [26]: # Generate input values between 0 and 1
x = torch.linspace(0.001, 1, 100) # Avoid log(0) by starting from 0.001


# Compute the Logarithm of the input values
y = torch.log(x)

# Create the plot
plt.figure(figsize=(4, 3))
plt.plot(x, y, label='log(x)')
plt.title('Logarithm of Input Values Between 0 and 1')
plt.xlabel('Input Value')
plt.ylabel('Logarithm')
plt.legend()
plt.grid(True)
plt.show()
```



**TODO:** 7c) Why is maximizing the likelihood equivalent to maximizing the log likelihood? (1 point)

**ANSWER:** YOUR ANSWER GOES HERE

Because the logarithm is a monotonically increasing function. This means that if you increase the likelihood, its logarithm also increases, and vice versa. 

Instead of maximizing the log likelihood, in practice we minimize the negative log likelihood, because a loss usually is a number in  $[0, \infty]$ , where small numbers are good. Finally, the negative log likelihood is simply the sum of the single logarithms due to

$$\log(a * b * c) = \log(a) + \log(b) + \log(c)$$

This is great because a single probability close to zero will not cause the whole loss to be tiny any more, we now use addition instead of multiplication!

**TODO:** 7d) Calculate the **negative log likelihood loss**, store it in `nll` and print it. Sometimes, we also use the **average negative log likelihood loss**, so average `nll` by the number of bigrams evaluated and print the average as well. **(3 points)**

**HINT:** You can start by copying the second last code cell above, where we calculated the probabilities of bigrams. Then apply `torch.log` to these probabilities and sum the negative logarithms.

```
In [27]: # YOUR CODE GOES HERE

nll = 0 # initialize negative log likelihood loss
num_bigrams = 0 # counter for the total number of bigrams

for w in words: # loop over each word in the dataset
    chs = ['.'] + list(w) + ['.'] # add start and end markers to each word
    for ch1, ch2 in zip(chs, chs[1:]): # iterate through each bigram in the word
        ix1 = stoi[ch1] # index for first character
        ix2 = stoi[ch2] # index for second character
        prob = P[ix1, ix2] # get the probability of the bigram
        nll += -torch.log(prob) # add the negative log probability to nll
        num_bigrams += 1 # increment the bigram count

# Average negative log likelihood
avg_nll = nll / num_bigrams

# Print results
print("Negative Log Likelihood Loss:", nll.item())
print("Average Negative Log Likelihood Loss:", avg_nll.item())
```

Negative Log Likelihood Loss: 559951.5625  
Average Negative Log Likelihood Loss: 2.4543561935424805



This summarizes our first bigram model, which simply counts occurrences of bigrams and generates new names based on the corresponding probability distribution. The average log likelihood loss of this model is around 2.45.

The following visualization includes the whole bigram model with negative log likelihood loss:



No description has been provided for this image

---

## 8. The Neural Network Approach

Let's try another approach: a neural network! Our neural network will still be a bigram language model: It receives a single character as an input, processes it using some weights or some parameters  $w$  and outputs the probability distribution over the next character. We will evaluate the parameters of the neural net using the negative log likelihood loss, which means comparing the output probability distribution and the label (the identity of the next

character). We are going to use gradient-based optimization to tune the parameters of this network to minimize the loss with the goal to better predict the next character. In the end, the result will look quite similar to the intuitive approach counting occurrences of bigrams above!

First, let's create the training set for the neural network made of two lists, the inputs and the targets:

```
In [28]: # create the training set of bigrams (x,y)
xs, ys = [], []

for w in words[:1]: # only use the first word for now, which contains 5 examples
    chs = ['.'] + list(w) + ['.']
    for ch1, ch2 in zip(chs, chs[1:]):
        ix1 = stoi[ch1]
        ix2 = stoi[ch2]
        print(ch1, ch2)
        xs.append(ix1) # input is the first character of the bigram
        ys.append(ix2) # target is the second character of the bigram

xs = torch.tensor(xs) # convert the list to a tensor
ys = torch.tensor(ys) # convert the list to a tensor

. e
e m
m m
m a
a .
```

We print the 5 resulting inputs and targets:

```
In [29]: xs
```

```
Out[29]: tensor([ 0,  5, 13, 13,  1])
```

```
In [30]: ys
```

```
Out[30]: tensor([ 5, 13, 13,  1,  0])
```

**TODO:** 8) The following code cell applies **one-hot-encoding**. Shortly explain what this is and why it is needed! **(2 points)**

**ANSWER:** YOUR ANSWER GOES HERE

One-hot encoding allows us to convert categorical data (like characters) into a numerical format that can be used as input for neural networks. Neural networks expect numerical input, and one-hot vectors ensure each character has a unique, distinguishable vector representation without introducing any inherent ordering or numerical relationships between characters.






```
In [31]: # apply one-hot encoding to the input
import torch.nn.functional as F
xenc = F.one_hot(xs, num_classes=27).float()
print(xenc)
print(xenc.shape)
```

```
tensor([[1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
        0., 0., 0., 0., 0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
        0., 0., 0., 0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0.,
        0., 0., 0., 0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0.,
        0., 0., 0., 0., 0., 0., 0., 0.],
        [0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
        0., 0., 0., 0., 0., 0., 0., 0.]])
torch.Size([5, 27])
```

```
In [32]: # we can also visualize the one-hot encoding like this:
plt.imshow(xenc)
```

```
Out[32]: <matplotlib.image.AxesImage at 0x1b37488fe60>
```



```
In [33]: # check the data type of the one-hot encoding
xenc.dtype
```

```
Out[33]: torch.float32
```

## 9. The First Neuron

A neuron consists of a simple dot product  $x \cdot W + b$ . We don't use a **bias**  $b$  here, so let's first initialize the **weights**  $W$  randomly sampled from a normal distribution, i.e., most weights will be around 0:

```
In [34]: W = torch.randn((27, 1))
W
```

**TODO:** 9) Apply the matrix multiplication (denoted by `@` in PyTorch) to get the **logits**. Store the result in a variable called `logits`. Which shape will the result have and why? Check your predicted shape by printing the result! **(2 points)**

```
In [35]: # YOUR CODE GOES HERE
```

```
logits = xenc @ W
print("Shape of xenc : ", xenc.shape)
print("Shape of W :", W.shape)
print("Shape of logits : ", logits.shape)
```

```
Shape of xenc : torch.Size([5, 27])
Shape of W : torch.Size([27, 1])
Shape of logits : torch.Size([5, 1])
```

**ANSWER:** YOUR ANSWER GOES HERE

When we multiply xenc (of shape (num\_bigrams, 27)) by W (of shape (27, 1)), the resulting shape will be (num\_bigrams, 1). Each row in logits will represent the output (or logit) for one bigram in the dataset.

Since our num\_bigrams size is 5, the result is of shape (5,1).

## 10. Creating 27 Neurons

Now we want 27 neurons instead of just one, because we eventually want to output the 27 probabilities for each character to be the next. So we want the weight matrix to be 27x27 instead of 27x1, and we will in parallel evaluate the 5 inputs on 27 neurons, so the output will be 5x27 (matrix multiplication of a 5x27 input with 27x27 weight matrix).

```
In [36]: # randomly initialize 27 neurons' weights. Each neuron receives 27 inputs
g = torch.Generator().manual_seed(2147483647)
W = torch.randn((27, 27), generator=g)
print(W)
```

```
tensor([[ 1.5674e+00, -2.3729e-01, -2.7385e-02, -1.1008e+00,  2.8588e-01,
         -2.9643e-02, -1.5471e+00,  6.0489e-01,  7.9136e-02,  9.0462e-01,
        -4.7125e-01,  7.8682e-01, -3.2843e-01, -4.3297e-01,  1.3729e+00,
         2.9334e+00,  1.5618e+00, -1.6261e+00,  6.7716e-01, -8.4039e-01,
         9.8488e-01, -1.4837e-01, -1.4795e+00,  4.4830e-01, -7.0730e-02,
         2.4968e+00,  2.4448e+00],
        [-6.7006e-01, -1.2199e+00,  3.0314e-01, -1.0725e+00,  7.2762e-01,
         5.1114e-02,  1.3095e+00, -8.0220e-01, -8.5042e-01, -1.8068e+00,
         1.2523e+00, -1.2256e+00,  1.2165e+00, -9.6478e-01, -2.3211e-01,
        -3.4762e-01,  3.3244e-01, -1.3263e+00,  1.1224e+00,  5.9641e-01,
         4.5846e-01,  5.4011e-02, -1.7400e+00,  1.1560e-01,  8.0319e-01,
         5.4108e-01, -1.1646e+00],
        [ 1.4756e-01, -1.0006e+00,  3.8012e-01,  4.7328e-01, -9.1027e-01,
        -7.8305e-01,  1.3506e-01, -2.1161e-01, -1.0406e+00, -1.5367e+00,
         9.3743e-01, -8.8303e-01,  1.7457e+00,  2.1346e+00, -8.5614e-01,
         5.4082e-01,  6.1690e-01,  1.5160e+00, -1.0447e+00, -6.6414e-01,
        -7.2390e-01,  1.7507e+00,  1.7530e-01,  9.9280e-01, -6.2787e-01,
         7.7023e-02, -1.1641e+00],
        [ 1.2473e+00, -2.7061e-01, -1.3635e+00,  1.3066e+00,  3.2307e-01,
         1.0358e+00, -8.6249e-01, -1.2575e+00,  9.4180e-01, -1.3257e+00,
        ...truncated...])
```

```
In [37]: xenc = F.one_hot(xs, num_classes=27).float() # input to the network: one-hot encoding
logits = xenc @ W # predict log-counts
```

```
print(logits)
tensor([[ 1.5674e+00, -2.3729e-01, -2.7385e-02, -1.1008e+00,  2.8588e-01,
         -2.9643e-02, -1.5471e+00,  6.0489e-01,  7.9136e-02,  9.0462e-01,
        -4.7125e-01,  7.8682e-01, -3.2843e-01, -4.3297e-01,  1.3729e+00,
         2.9334e+00,  1.5618e+00, -1.6261e+00,  6.7716e-01, -8.4039e-01,
         9.8488e-01, -1.4837e-01, -1.4795e+00,  4.4830e-01, -7.0730e-02,
         2.4968e+00,  2.4448e+00],
       [ 4.7236e-01,  1.4830e+00,  3.1748e-01,  1.0588e+00,  2.3982e+00,
         4.6827e-01, -6.5650e-01,  6.1662e-01, -6.2197e-01,  5.1007e-01,
         1.3563e+00,  2.3445e-01, -4.5585e-01, -1.3132e-03, -5.1161e-01,
         5.5570e-01,  4.7458e-01, -1.3867e+00,  1.6229e+00,  1.7197e-01,
         9.8846e-01,  5.0657e-01,  1.0198e+00, -1.9062e+00, -4.2753e-01,
        -2.1259e+00,  9.6041e-01],
       [ 1.9359e-01,  1.0532e+00,  6.3393e-01,  2.5786e-01,  9.6408e-01,
        -2.4855e-01,  2.4756e-02, -3.0404e-02,  1.5622e+00, -4.4852e-01,
        -1.2345e+00,  1.1220e+00, -6.7381e-01,  3.7882e-02, -5.5881e-01,
        -8.2709e-01,  8.2253e-01, -7.5100e-01,  9.2778e-01, -1.4849e+00,
        -2.1293e-01, -1.1860e+00, -6.6092e-01, -2.3348e-01,  1.5447e+00,
         6.0061e-01, -7.0909e-01],
       [ 1.9359e-01,  1.0532e+00,  6.3393e-01,  2.5786e-01,  9.6408e-01,
        -2.4855e-01,  2.4756e-02, -3.0404e-02,  1.5622e+00, -4.4852e-01,
        ...truncated...
```

These numbers tell us the firing rate of each of the 27 neurons on the 5 inputs. For example, the result at row 3 and column 13 is giving us the firing rate of the 13th neuron looking at the third input:

```
In [38]: (xenc @ W)[3,13]
```


```
Out[38]: tensor(0.0379)
```

This is equivalent to a dot product between the third input and the 13th column of  $W$ , so using matrix multiplication we can very efficiently evaluate the dot product between lots of input examples in a batch and lots of neurons:

```
In [39]: (xenc[3] * W[:,13]).sum() # element-wise multiplication and sum is equivalent to ma
```

```
Out[39]: tensor(0.0379)
```

Here we can see the steps so far, including one-hot-encoding, weights and logit calculation for input 'e':



No description has been provided for this image

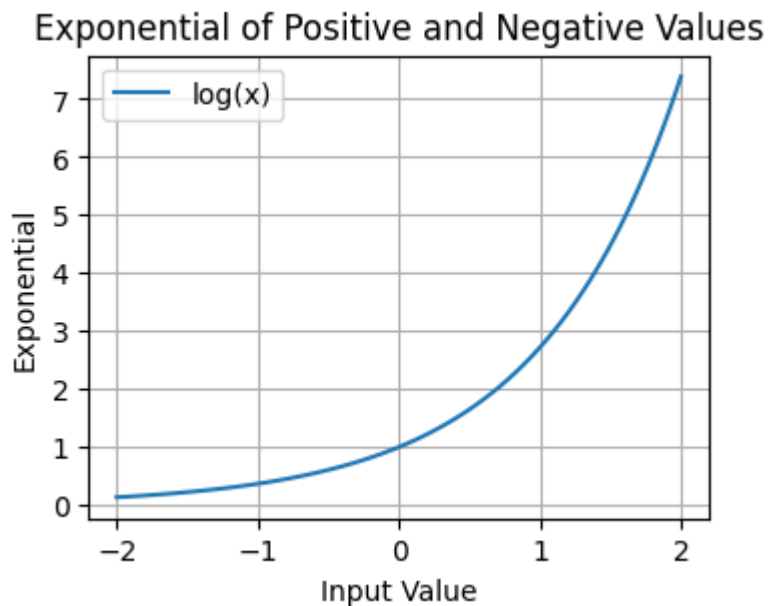
Note that in the image we only plot a single input 'e', whereas in the code we process 5 inputs in parallel.

We see that the logits are positive and negative. They can be interpreted as **log counts** (=logarithms of the counts of each bigram). We can elementwise exponentiate to transform them into numbers in  $(0, \infty)$ , where negative numbers go to  $(0, 1)$  and positive numbers go to  $(1, \infty)$ :

```
In [40]: # Generate input values between -2 and 2
x = torch.linspace(-2, 2, 100) #

# Compute the exponential of the input values
y = torch.exp(x)

# Create the plot
plt.figure(figsize=(4, 3))
plt.plot(x, y, label='log(x)')
plt.title('Exponential of Positive and Negative Values')
plt.xlabel('Input Value')
plt.ylabel('Exponential')
plt.legend()
plt.grid(True)
```



**TODO:** 10a) Calculate the counts from the log counts by applying exponential function. Store the result in `counts` and print it. Normalize the counts like above, store the result in `probs` and print it. **(2 points)**

```
In [41]: # YOUR CODE GOES HERE


# Calculate counts from log counts
counts = torch.exp(logits) # Apply exponential function to each element in logits
print("Counts:\n", counts)

# Normalize counts to get probabilities
probs = counts / counts.sum(dim=1, keepdims=True) # Sum across rows and divide
print("Probabilities:\n", probs)
```

Counts:

```
tensor([[ 4.7940,  0.7888,  0.9730,  0.3326,  1.3309,  0.9708,  0.2129,  1.8311,
          1.0824,  2.4710,  0.6242,  2.1964,  0.7200,  0.6486,  3.9469, 18.7908,
          4.7673,  0.1967,  1.9683,  0.4315,  2.6775,  0.8621,  0.2277,  1.5656,
          0.9317, 12.1434, 11.5281],
        [ 1.6038,  4.4060,  1.3737,  2.8830, 11.0032,  1.5972,  0.5187,  1.8527,
          0.5369,  1.6654,  3.8818,  1.2642,  0.6339,  0.9987,  0.5995,  1.7432,
          1.6073,  0.2499,  5.0680,  1.1876,  2.6871,  1.6596,  2.7728,  0.1486,
          0.6521,  0.1193,  2.6128],
        [ 1.2136,  2.8669,  1.8850,  1.2942,  2.6224,  0.7799,  1.0251,  0.9701,
          4.7691,  0.6386,  0.2910,  3.0710,  0.5098,  1.0386,  0.5719,  0.4373,
          2.2763,  0.4719,  2.5289,  0.2265,  0.8082,  0.3054,  0.5164,  0.7918,
          4.6866,  1.8232,  0.4921],
        [ 1.2136,  2.8669,  1.8850,  1.2942,  2.6224,  0.7799,  1.0251,  0.9701,
          4.7691,  0.6386,  0.2910,  3.0710,  0.5098,  1.0386,  0.5719,  0.4373,
          2.2763,  0.4719,  2.5289,  0.2265,  0.8082,  0.3054,  0.5164,  0.7918,
          4.6866,  1.8232,  0.4921],
        [ 0.5117,  0.2953,  1.3541,  0.3422,  2.0701,  1.0524,  3.7043,  0.4483,
          0.4272,  0.1642,  3.4984,  0.2936,  3.3753,  0.3811,  0.7929,  0.7064,
          1.3944,  0.2655,  3.0723,  1.8156,  1.5816,  1.0555,  0.1755,  1.1225,
          ...truncated...])
```

Here are the visualized steps including exponentiation and normalization:

No description has been provided for this image

**TODO:** 10b) Do the last two maths operations in the cell above look familiar? How is this transformation called? **(1 point)**

**ANSWER:** YOUR ANSWER GOES HERE

Yes, the last two math operations combined are known as the softmax transformation.

Let's take a look at the resulting "probabilities" for the first character (of course they are meaningless because the network is still untrained, the weights are randomly initialized):

```
In [42]: probs[0]
```

```
In [43]: probs[0].shape
```

```
Out[43]: torch.Size([27])
```

```
In [44]: probs.shape
```

```
Out[44]: torch.Size([5, 27])
```

## 11. Summary (so far)

Let's summarize the code so far to get an overview of the steps:

```
In [45]: xs # input to the network ('.emma')
```

```
Out[45]: tensor([ 0,  5, 13, 13,  1])
```

```
In [46]: ys # target for the network ('emma.')
```

```
Out[46]: tensor([ 5, 13, 13,  1,  0])
```

```
In [47]: # randomly initialize 27 neurons' weights. each neuron receives 27 inputs
g = torch.Generator().manual_seed(2147483647)
W = torch.randn((27, 27), generator=g)
```

```
In [48]: xenc = F.one_hot(xs, num_classes=27).float() # input to the network: one-hot encoding
logits = xenc @ W # predict log-counts
counts = logits.exp() # counts, equivalent to N
probs = counts / counts.sum(1, keepdims=True) # probabilities for next character
```

```
In [49]: probs.shape # output probabilities of shape 5x27
```

```
Out[49]: torch.Size([5, 27])
```

```
In [50]: # compute the negative log likelihood loss for the first 5 bigrams
nlls = torch.zeros(5)
for i in range(5):
    # i-th bigram:
    x = xs[i].item() # input character index
    y = ys[i].item() # label character index
    print('-----')
    print(f'bigram example {i+1}: {itos[x]}{itos[y]} (indices {x},{y})')
    print('input to the neural net:', x)
    print('output probabilities from the neural net:', probs[i])
    print('label (actual next character):', y)
    p = probs[i, y]
    print('probability assigned by the net to the the correct character:', p.item())
    logp = torch.log(p)
    print('log likelihood:', logp.item())
    nll = -logp
    print('negative log likelihood:', nll.item())
    nlls[i] = nll


print('=====')
print('average negative log likelihood, i.e. loss =', nlls.mean().item())
```

```

-----
bigram example 1: .e (indices 0,5)
input to the neural net: 0
output probabilities from the neural net: tensor([0.0607, 0.0100, 0.0123, 0.0042, 0.
0168, 0.0123, 0.0027, 0.0232, 0.0137,
          0.0313, 0.0079, 0.0278, 0.0091, 0.0082, 0.0500, 0.2378, 0.0603, 0.0025,
          0.0249, 0.0055, 0.0339, 0.0109, 0.0029, 0.0198, 0.0118, 0.1537, 0.1459])
label (actual next character): 5
probability assigned by the net to the the correct character: 0.01228625513613224
log likelihood: -4.399273872375488
negative log likelihood: 4.399273872375488
-----
bigram example 2: em (indices 5,13)
input to the neural net: 5
output probabilities from the neural net: tensor([0.0290, 0.0796, 0.0248, 0.0521, 0.
1989, 0.0289, 0.0094, 0.0335, 0.0097,
          0.0301, 0.0702, 0.0228, 0.0115, 0.0181, 0.0108, 0.0315, 0.0291, 0.0045,
          0.0916, 0.0215, 0.0486, 0.0300, 0.0501, 0.0027, 0.0118, 0.0022, 0.0472])
label (actual next character): 13
probability assigned by the net to the the correct character: 0.018050700426101685
log likelihood: -4.014570713043213
negative log likelihood: 4.014570713043213
...truncated...

```

Here is a visualization of the first feedforward pass of input 'e' through the neuron, including negative log likelihood loss calculation. You can see in the output above that the loss is roughly 4 (see 'bigram example 2' output).

 No description has been provided for this image

We see from the output above that the average loss is 3.7, which is quite high. We can change the random seed for sampling  $W$  for randomly changing the loss. But of course, we can do better: The loss calculation is only made up of differentiable operations (multiplication, addition, exponential, division...) We can minimize the loss by computing the gradients of the loss with respect to  $W$ .

## 12. Optimization

To calculate the loss, we basically need to pluck out the predicted probabilities at the indices of the target characters. For example, for the first bigram, the input character is '.' (index 0) and the target character is 'e' (index 5), and the predicted probability is stored in `probs[0,5]`.

```

In [51]: # we need the following probabilities for calculating the loss:
         probs[0,5], probs[1,13], probs[2,13], probs[3,1], probs[4,0]

```

```
Out[51]: (tensor(0.0123),
          tensor(0.0181),
          tensor(0.0267),
          tensor(0.0737),
          tensor(0.0150))
```

This is equivalent to:

```
In [52]: probs[torch.arange(5), ys]
```

```
Out[52]: tensor([0.0123, 0.0181, 0.0267, 0.0737, 0.0150])
```

We can use this in the loss calculation at the end of the forward pass:

```
In [53]: # randomly initialize 27 neurons' weights. each neuron receives 27 inputs
g = torch.Generator().manual_seed(2147483647)
W = torch.randn((27, 27), generator=g, requires_grad=True) # we need to compute gra
```

```
In [54]: # forward pass
xenc = F.one_hot(xs, num_classes=27).float() # input to the network: one-hot encodi
logits = xenc @ W # predict log-counts
counts = logits.exp() # counts, equivalent to N
probs = counts / counts.sum(1, keepdims=True) # probabilities for next character
loss = -probs[torch.arange(5), ys].log().mean() # average negative log likelihood L
```

```
In [55]: print(loss.item()) # print the loss
```

```
3.7693049907684326
```

Note that the loss is the same as above. Finally we can do the backward pass:

```
In [56]: # backward pass
W.grad = None # set to zero the gradient (None is a special value that PyTorch reco
loss.backward() # compute the gradients
```

Something magical happened when `loss.backward()` was run: PyTorch keeps track of all the operations in the forward pass, under the hood it builds a full computational graph, so it knows all the dependencies and all the mathematical operations inside the neural network. When we calculate the loss and call a `.backward()` on it, PyTorch fills in the partial derivatives of all the intermediate steps. So now we can look at the gradient and see that it is not zero anymore:

```
In [57]: W.grad
```

```
In [58]: W.shape, W.grad.shape
```

```
Out[58]: (torch.Size([27, 27]), torch.Size([27, 27]))
```

Both `W` and `W.grad` are of shape 27x27, and each entry of `W.grad` is telling us the influence of that weight on the loss function. For example, since `W.grad[0,0]` is positive,



slightly increasing `W[0,0]` will lead to a slightly bigger loss. Therefore we use the *negative* gradient for updating the weights:

```
In [59]: W.data += -0.1 * W.grad # update the weights with Learning rate 0.1
```

Let's check that the loss has really decreased using this one **gradient descent** step:

```
In [60]: # forward pass
xenc = F.one_hot(xs, num_classes=27).float() # input to the network: one-hot encodi
logits = xenc @ W # predict log-counts
counts = logits.exp() # counts, equivalent to N
probs = counts / counts.sum(1, keepdims=True) # probabilities for next character
loss = -probs[torch.arange(5), ys].log().mean()
print(loss.item())
```

3.7492127418518066

## 13. Optimizing over the Whole Dataset

Now let's apply gradient descent to the whole dataset:

```
In [61]: # create the dataset
xs, ys = [], []
for w in words:
    chs = ['.'] + list(w) + ['.']
    for ch1, ch2 in zip(chs, chs[1:]):
        ix1 = stoi[ch1]
        ix2 = stoi[ch2]
        xs.append(ix1)
        ys.append(ix2)
xs = torch.tensor(xs)
ys = torch.tensor(ys)
num = xs.nelement() # number of elements in the tensor xs (5 in '.emma.', here more
print('number of examples: ', num)

# initialize the 'network'
g = torch.Generator().manual_seed(2147483647)
W = torch.randn((27, 27), generator=g, requires_grad=True)
```

number of examples: 228146

```
In [62]: # gradient descent
for k in range(500):

    # forward pass
    xenc = F.one_hot(xs, num_classes=27).float() # input to the network: one-hot en
    logits = xenc @ W # predict log-counts - one-hot input will pluck the correct r
    counts = logits.exp() # counts, equivalent to N
    probs = counts / counts.sum(1, keepdims=True) # probabilities for next characte
    loss = -probs[torch.arange(num), ys].log().mean() + 0.01*(W**2).mean()
    print(loss.item())

    # backward pass
```

```
W.grad = None # set to zero the gradient
loss.backward()

# update
W.data += -50 * W.grad # even a large learning rate is fine here, because the b
```

```
3.7686190605163574
3.3788068294525146
3.161090850830078
3.027186155319214
2.9344842433929443
2.8672313690185547
2.816654682159424
2.777146577835083
2.745253801345825
2.7188303470611572
2.696505546569824
2.6773719787597656
2.6608052253723145
2.6463515758514404
2.633665084838867
2.622471570968628
2.6125476360321045
2.6037068367004395
2.595794916152954
2.5886809825897217
...truncated...
```

We see that the loss is converging roughly to 2.48, which is similar to the loss we achieved all the way up in the original bigram model without a neural network (2.45). Back then, we achieved this loss by counting, now by optimizing weights. That makes sense because fundamentally we're not using any additional information, we're still just taking in the previous character and trying to predict the next one, but instead of doing it explicitly by counting and normalizing we are doing it with gradient-based learning. The explicit approach happens to very well optimize the loss function without any need for a gradient based optimization because the setup for bigram language models is so straightforward, we can just afford to estimate those probabilities directly and maintain them in a table. But the gradient-based approach is significantly more flexible, so we've actually gained a lot because we can expand this approach and complexify the neural net.

**TODO (optional):** Remember the model smoothing used in the first bigram model by adding a number to the counts to avoid zero probabilities. Can you think of a smoothing equivalent in gradient descent based bigram model?

**HINT:** Think of the weight matrix  $W$  and its influence on the result!

**ANSWER:** YOUR ANSWER GOES HERE

The code above already includes a smoothing:

```
loss = -probs[torch.arange(num), ys].log().mean() + 0.01*(W**2).mean()
```

While the last part `+ 0.01*(W**2).mean()` is a regularization term which is proportional to `W**2` which penalizes large values in `W` and encourages a smoother and less extreme distribution of weights.



Let's see how the weights and the probabilities have changed after training:

```
In [63]: print(W)
```

```
tensor([[ -3.2895e+00,  1.9656e+00,  7.4837e-01,  9.1425e-01,  1.0058e+00,
          9.0709e-01, -3.8237e-01,  8.2961e-02,  3.4820e-01, -3.9621e-02,
          1.3657e+00,  1.5675e+00,  9.3350e-01,  1.4125e+00,  6.1800e-01,
         -4.3777e-01, -1.7532e-01, -1.7782e+00,  9.7520e-01,  1.2013e+00,
          7.4990e-01, -1.9133e+00, -4.8333e-01, -6.7977e-01, -1.4532e+00,
         -1.3781e-01,  4.0890e-01],
        [ 2.2965e+00, -1.8025e-01, -2.0725e-01, -3.4589e-01,  4.4363e-01,
          3.6399e-02, -1.5385e+00, -1.3327e+00,  1.2488e+00,  9.0269e-01,
         -1.2949e+00, -1.5915e-01,  1.3296e+00,  8.9294e-01,  2.0966e+00,
         -2.1706e+00, -1.9615e+00, -2.2076e+00,  1.5854e+00,  5.1385e-01,
          2.9205e-02, -5.5165e-01,  2.2184e-01, -1.3719e+00, -1.2585e+00,
          1.1198e+00, -4.2190e-01],
        [ 1.3027e+00,  2.3723e+00,  2.0947e-01, -1.0908e+00,  7.2807e-01,
          3.1045e+00, -1.1616e+00, -1.1911e+00,  2.6474e-01,  1.9679e+00,
         -1.0697e+00, -1.2726e+00,  1.1984e+00, -1.0892e+00, -1.0602e+00,
          1.2181e+00, -1.1316e+00, -1.0999e+00,  3.3608e+00, -8.4553e-01,
         -1.1455e+00,  3.7126e-01, -1.1586e+00, -1.1152e+00, -1.2375e+00,
          9.7738e-01, -1.3179e+00],
        [ 6.1551e-01,  2.7858e+00, -1.6589e+00, -1.8775e-01, -1.4958e+00,
          2.3871e+00, -1.6076e+00, -1.5478e+00,  2.5773e+00,  1.6614e+00,
          ...truncated...
```

```
In [64]: # compute the negative log likelihood loss for the first 5 bigrams
nlls = torch.zeros(5)
for i in range(5):
    # i-th bigram:
    x = xs[i].item() # input character index
    y = ys[i].item() # label character index
    print('-----')
    print(f'bigram example {i+1}: {itos[x]}{itos[y]} (indices {x},{y})')
    print('input to the neural net:', x)
    print('output probabilities from the neural net:', probs[i])
    print('label (actual next character):', y)
    p = probs[i, y]
    print('probability assigned by the net to the the correct character:', p.item())
    logp = torch.log(p)
    print('log likelihood:', logp.item())
    nll = -logp
    print('negative log likelihood:', nll.item())
    nlls[i] = nll


print('=====')
print('average negative log likelihood, i.e. loss =', nlls.mean().item())
```

```

-----
bigram example 1: .e (indices 0,5)
input to the neural net: 0
output probabilities from the neural net: tensor([0.0007, 0.1373, 0.0407, 0.0480, 0.
0526, 0.0476, 0.0131, 0.0209, 0.0272,
          0.0185, 0.0754, 0.0922, 0.0489, 0.0790, 0.0357, 0.0124, 0.0161, 0.0032,
          0.0510, 0.0639, 0.0407, 0.0028, 0.0119, 0.0097, 0.0045, 0.0168, 0.0290],
          grad_fn=<SelectBackward0>)
label (actual next character): 5
probability assigned by the net to the the correct character: 0.04764774069190025
log likelihood: -3.0439200401306152
negative log likelihood: 3.0439200401306152
-----
bigram example 2: em (indices 5,13)
input to the neural net: 5
output probabilities from the neural net: tensor([0.1943, 0.0330, 0.0062, 0.0077, 0.
0188, 0.0618, 0.0044, 0.0064, 0.0077,
          0.0398, 0.0032, 0.0089, 0.1583, 0.0374, 0.1303, 0.0132, 0.0045, 0.0014,
          0.0953, 0.0419, 0.0282, 0.0038, 0.0226, 0.0030, 0.0067, 0.0520, 0.0090],
          grad_fn=<SelectBackward0>)
label (actual next character): 13
probability assigned by the net to the the correct character: 0.037398409098386765
...truncated...

```

**TODO:** 13) Here is the final visualization after training. Can you interpret the weights, have they improved? Can you compare the trained weights to the weights from our first approach using simple counting instead of a neuron? **(2 points)**

 No description has been provided for this image

**ANSWER:** YOUR ANSWER GOES HERE

The weights have improved, as indicated by the decrease in the average negative log likelihood from 3.7 to 2.5. This shows that the model has become better at predicting the correct characters, with higher probabilities assigned to them after training.

The weights from the first approach using simple counting can be thought of as counts of occurrences normalized by the total occurrences of the first character in each bigram. However, the weights of the trained neuron are an improvement in the following way:

The weights do not just reflect direct counts but learned patterns that capture dependencies between characters more effectively than raw counts. Means, the weight is higher if a dependency occurred often. Also the trained neuron reflects a smoothing effect: even if a certain bigram wasn't seen during the training, the model still assign a nonzero probability to it.



## 14. Sampling from the Neural Net

```
In [65]: # finally, sample from the 'neural net' model
g = torch.Generator().manual_seed(2147483647)
```

```

for i in range(20):

    out = []
    ix = 0
    while True:

        # -----
        # BEFORE:
        #p = P[ix] # bigram probabilities by counting
        # -----
        # NOW:
        xenc = F.one_hot(torch.tensor([ix]), num_classes=27).float()
        logits = xenc @ W # predict log-counts
        counts = logits.exp() # counts, comparable to N
        p = counts / counts.sum(1, keepdims=True) # probabilities for next character
        # -----

        ix = torch.multinomial(p, num_samples=1, replacement=True, generator=g).item()
        out.append(itos[ix])
        if ix == 0:
            break
    print(''.join(out))

```

```

junide.
janasah.
p.
cfay.
a.
nn.
kohin.
tolian.
juwe.
ksahnaauranilevias.
dedainrwieta.
ssonielylarte.
faveumerifontume.
phynslenaruani.
core.
yaenon.
ka.
jabdinerimikimaynin.
anaasn.
ssorionsush.
...truncated...

```

We get nearly identical samples as in the first model!

## 15. Conclusion

To summarize, we introduced the bigram character level language model, we saw how we can train the model, how we can sample from the model and how we can evaluate the quality of the model using the negative log likelihood loss. We trained the model in two completely different ways that actually get the same result and the same model: First, we just

counted the frequency of all the bigrams and normalized. Second, we used the negative log likelihood loss as a guide to optimizing the counts matrix so that the loss was minimized in a gradient-based framework. Both approaches gave the same result, but the gradient-based framework is much more flexible and we can extend it to more complex settings.

```
In [66]: import sys
import os

# Add the parent directory to the sys.path
sys.path.append(os.path.abspath(os.path.join('../')))

# truncate long cell output to avoid large pdf files
from helpers.truncate_output import truncate_long_notebook_output
truncated = truncate_long_notebook_output('1_Bigram_Language_Model__Gruppe2.ipynb')

# convert to pdf with nbconvert
if truncated:
    !jupyter nbconvert --to webpdf --allow-chromium-download TRUNCATED_1_Bigram_Lan
else:
    !jupyter nbconvert --to webpdf --allow-chromium-download 1_Bigram_Language_Model
```

Output in 1\_Bigram\_Language\_Model\_\_Gruppe2.ipynb not above threshold and so no new version made.

```
[NbConvertApp] Converting notebook 1_Bigram_Language_Model__Gruppe2.ipynb to webpdf
[NbConvertApp] WARNING | Alternative text is missing on 7 image(s).
[NbConvertApp] Building PDF
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 237739 bytes to 1_Bigram_Language_Model__Gruppe2.pdf
```

In [ ]:

